

COMP 650 “Thesis” Presentation

RAM and IJRR/RSS TMLKit Papers

Neil T. Dantam et. al., Rice University, 2016

Presented by Cannon Lewis

Images from original papers

Identified Problems

- Task and Motion Planning (TMP) presents challenges in completeness and scalability; existing approaches sacrifice one or the other.
- Existing TMP solutions are specific to certain classes of problems.
- Task planners and motion planners operate on very different principles, so achieving interaction between these layers is difficult.
- Motion planning for complex manipulation tasks requires a non-static configuration space representation.



Specific Contributions

- Useful abstractions of the task planning, motion planning, and TMP problems.
- TMKit, a modular system for performing interleaved TMP.
- The IDTMP algorithm for efficiently combining task planning and motion planning.
- The Amino scene-graph representation which allows for efficient configuration space updates.



Background - Task Planning

Definition 1. Task Language. *The task language is a set of strings of actions, defined by $\mathcal{L} = (\mathcal{P}, \mathcal{A}, \mathcal{E}, s^{[0]}, \mathcal{G})$, where,*

- \mathcal{P} is the state space ranging over variables p_0, \dots, p_n ,
- \mathcal{A} is the set of task operators, i.e., terminal symbols,
- $\mathcal{E} \subseteq (\mathbb{P}(\mathcal{P}) \times \mathcal{A} \times \mathbb{P}(\mathcal{P}))$ is the set of symbolic transitions. Each $e_i \in \mathcal{E}$ denotes transitions $\text{pre}(a_i) \xrightarrow{a_i} \text{eff}(a_i)$, where $\text{pre}(a_i) \subseteq \mathcal{P}$ is the precondition set, $a_i \in \mathcal{A}$ is the operator, and $\text{eff}(a_i) \subseteq \mathcal{P}$ is the effect set. We represent a concrete transition on a_i at step k from state $s^{[k]}$ to $s^{[k+1]}$ as $s^{[k+1]} = a_i(s^{[k]})$, where $s^{[k]} \in \text{pre}(a_i)$, $s^{[k+1]} \in \text{eff}(a_i)$, and for all state variables p_j independent of (i.e., free in) $\text{eff}(a_i)$, $p_j^{[k+1]} = p_j^{[k]}$,
- $s^{[0]} \in \mathcal{P}$ is the start state,
- $\mathcal{G} \subseteq \mathcal{P}$ is the set of accept states, i.e., the task goal.

Definition 2. Task Plan. *A task plan \mathbf{A} is a string in the task language \mathcal{L} , i.e., $\mathbf{A} \in \mathcal{L}$, where $\mathbf{A} = (a^{[0]}, a^{[1]}, \dots, a^{[h]})$, $a^{[k]} \in \mathcal{A}$, $s^{[k]} \in \text{pre}(a^{[k]})$, $s^{[k+1]} = a^{[k]}(s^{[k]})$, and $s^{[h]} \in \mathcal{G}$.*

- Domain consists of discrete representation of task to be solved.
- Plan consists of discrete actions to transition from start state to goal state.



Background - Task Planning

- We would like to be able to generate alternate task plans, so that motion planner feedback can influence task plan choice.
- We would also like to be able to reuse work from previous planning rounds.
- Dominant approaches in task planning are heuristic search (e.g. A*) and constraint based methods (e.g. SMT).



Background - Task Planning

- Satisfiability Modulo Theories (SMT) is essentially SAT solving with additional rules (“theories”) for defining constraints in non-boolean domains.
- SMT effectively provides a high-level interface for SAT solving.
- Incremental solving with Z3 (pushing/popping constraints).
- Want to encode failed motion plans with constraints.

The logo for the Z3 theorem prover, featuring the letters 'Z3' in a large, blue, 3D-style font with a white outline and a slight shadow.

Background - Task Languages

- Many task language notations exist.
- Most popular are PDDL, LTL, or special purpose formal languages.
- All notations roughly equivalent in expressibility.
- Ideal TMP solution is notation agnostic.
- In experiments for this paper, PDDL is used.

```
(define (domain transfer-linear)
  (:types block location - object)
  (:functions (position ?obj - block) - location)
  (:derived (occupied ?loc - location)
    (exists (?obj - block) (= (position ?obj)
                              ?loc)))


  (:action transfer
    :parameters (?obj - block
                  ?dst - location)
    :precondition (not (occupied ?dst))
    :effect (and (= (position ?obj)
                    ?dst))))
```

Background - Motion Planning

- Robots typically modeled as kinematic trees.
- For task and motion planning, need to be able to modify kinematic tree as objects are manipulated.
- Sampling based planners are the dominant method for high DoF motion planning.
- Probabilistic completeness a consideration for TMP.
- For this work, RRT-Connect is used.



Background - Motion Planning

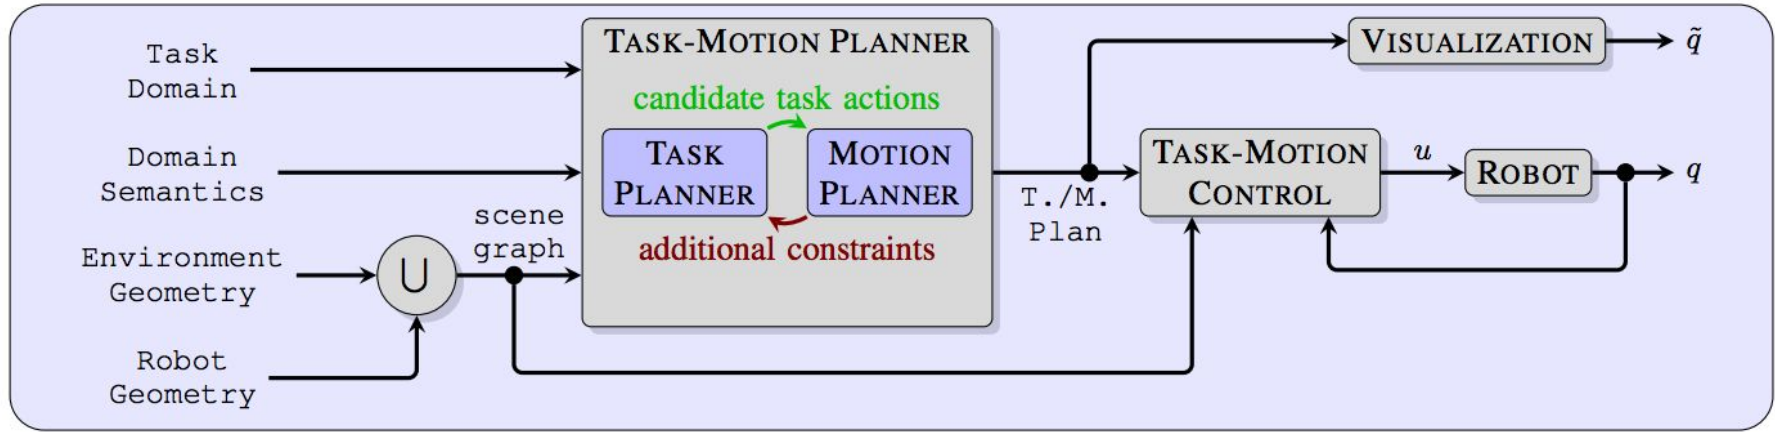


Background - Task and Motion Planning

- Seek to establish correspondence between task operators and motion planning problems.
- Due to probabilistic nature of sampling planners, cannot prove nonexistence of motion plans for operators.
- Task and motion plan consists of paired task operators and motion plans.
- Task and motion planning steps are typically interleaved, failed motion plans revisited.



Background - Task and Motion Planning



Related Work

- Many existing methods for TMP, mostly for specific domains
- Existing methods typically focus on performance.
- TMP methods include semantic attachments, knowledge bases, Hierarchical Planning in the Now, discretized motion planning.
- aSyMov and Synergistic Framework ensure probabilistic completeness, but use different algorithms for performing task planning, motion planning, and at the TMP interface level than TMKit.



Related Work

- Some related TMP works deal directly with differential dynamics, though TMKit does not.
- Many other TMP works use roadmaps, which then avoid much of the repeated computation necessary in TMKit.
- However, roadmap methods decrease flexibility to changes in configuration space.
- A few TMP approaches require or make use of backtracking in lieu of establishing alternate task plans.



Contributions - Abstractions

Definition 1. Task Language. *The task language is a set of strings of actions, defined by $\mathcal{L} = (\mathcal{P}, \mathcal{A}, \mathcal{E}, s^{[0]}, \mathcal{G})$, where,*

- \mathcal{P} is the state space ranging over variables p_0, \dots, p_n ,
- \mathcal{A} is the set of task operators, i.e., terminal symbols,
- $\mathcal{E} \subseteq (\mathbb{P}(\mathcal{P}) \times \mathcal{A} \times \mathbb{P}(\mathcal{P}))$ is the set of symbolic transitions. Each $e_i \in \mathcal{E}$ denotes transitions $\text{pre}(a_i) \xrightarrow{a_i} \text{eff}(a_i)$, where $\text{pre}(a_i) \subseteq \mathcal{P}$ is the precondition set, $a_i \in \mathcal{A}$ is the operator, and $\text{eff}(a_i) \subseteq \mathcal{P}$ is the effect set. We represent a concrete transition on a_i at step k from state $s^{[k]}$ to $s^{[k+1]}$ as $s^{[k+1]} = a_i(s^{[k]})$, where $s^{[k]} \in \text{pre}(a_i)$, $s^{[k+1]} \in \text{eff}(a_i)$, and for all state variables p_j independent of (i.e., free in) $\text{eff}(a_i)$, $p_j^{[k+1]} = p_j^{[k]}$,
- $s^{[0]} \in \mathcal{P}$ is the start state,
- $\mathcal{G} \subseteq \mathcal{P}$ is the set of accept states, i.e., the task goal.

Definition 2. Task Plan. *A task plan \mathbf{A} is a string in the task language \mathcal{L} , i.e., $\mathbf{A} \in \mathcal{L}$, where $\mathbf{A} = (a^{[0]}, a^{[1]}, \dots, a^{[h]})$, $a^{[k]} \in \mathcal{A}$, $s^{[k]} \in \text{pre}(a^{[k]})$, $s^{[k+1]} = a^{[k]}(s^{[k]})$, and $s^{[h]} \in \mathcal{G}$.*

Contributions - Abstractions

Definition 3. Scene Graph. $\gamma = (\mathcal{Q}, \mathcal{L}, \mathcal{F})$, where

- $\mathcal{Q} \subseteq \mathbb{R}^n$ is a space of configurations,
- \mathcal{L} is a finite set of unique frame labels,
- \mathcal{F} is a finite set of kinematic frames (graph nodes), such that each frame $f_\ell = (\ell, \varrho_\ell, \varsigma_\ell, \mu_\ell)$, where,
 - $\ell \in \mathcal{L}$ is the unique label of frame of f_ℓ
 - $\varrho_\ell \in \mathcal{L}$ is the label of the parent of frame of f_ℓ , indicating graph edge connections
 - $\varsigma_\ell : \mathcal{Q} \mapsto \mathcal{SE}(3)$, maps from the configuration space to the workspace pose of f_ℓ relative to its parent ϱ_ℓ , indicating graph edge values
 - μ_ℓ is a rigid body mesh representing geometry attached to f_ℓ .

Definition 4. Motion Plan. A motion plan is a sequence of neighboring configurations $\mathbf{Q} = (q^{[0]}, q^{[1]}, \dots, q^{[n]})$ such that each $q^{[k]} \in \mathcal{Q}$ and $\|q^{[i+1]} - q^{[i]}\| < \epsilon_q$, for some small ϵ_q . The initial configuration is $\text{first}(\mathbf{Q}) = q^{[0]}$, and the final configuration is $\text{last}(\mathbf{Q}) = q^{[n]}$.



Contributions - Abstractions

Definition 5. Task-Motion Domain.

$$\mathfrak{D} = (\mathcal{L}, \sigma^{[0]}, \lambda_\alpha, \lambda_\rho, \Omega)$$

- \mathcal{L} is the task language, where $\mathcal{P} = \mathcal{P}_m \times \mathcal{P}_t$ represents the motion component \mathcal{P}_m and non-motion component \mathcal{P}_t of task state,
- $\sigma^{[0]} = (s^{[0]}, \gamma^{[0]}, q^{[0]})$ is initial the task-motion state: task state $s^{[0]}$, scene graph $\gamma^{[0]}$ and configuration $q^{[0]}$,
- $\lambda_\alpha : \Gamma \mapsto \mathcal{P}_m$ domain semantics to abstract the scene graph $\gamma \in \Gamma$ to the motion component task state $s_m \in \mathcal{P}_m$,
- $\lambda_\rho : \Gamma \times \mathcal{A} \mapsto \mathbb{P}(\mathcal{Q}) \times \Gamma$ domain semantics to refine the initial scene graph $\gamma^{[k]} \in \Gamma$ and task operator $a^{[k]} \in \mathcal{A}$ to a motion planning goal (a set of configurations) $\Theta^{[k]} \subseteq \mathcal{Q}$ for the action and a final scene graph $\gamma^{[k+1]} \in \Gamma$ via reparenting frames,
- $\Omega \subseteq \Gamma \times \mathcal{P}$ is the goal condition.

Definition 6. Task and Motion Plan. A task and motion plan is a sequence of task operators and motion plans, $\mathbf{T} = ((a^{[0]}, \mathbf{Q}^{[0]}), \dots, (a^{[h]}, \mathbf{Q}^{[h]}))$ where $(a^{[0]}, \dots, a^{[h]}) \in \mathcal{L}$ and for each step k , $(\Theta^{[k]}, \gamma^{[k+1]}) = \lambda_\rho(a^{[k]}, \gamma^{[k]})$ such that $\text{last}(\mathbf{Q}^{[k]}) \in \Theta^{[k]} \wedge \text{first}(\mathbf{Q}^{[k+1]}) = \text{last}(\mathbf{Q}^{[k]})$.

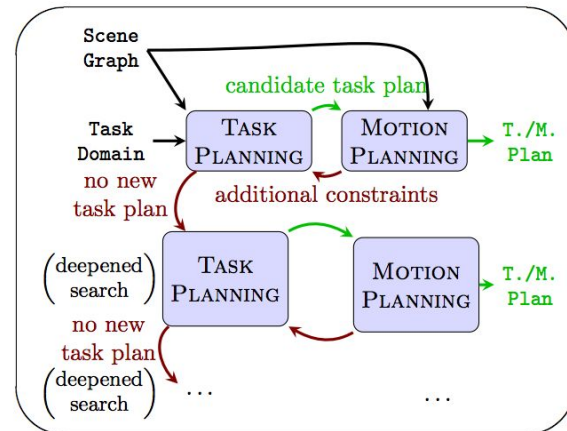
Contributions - Abstractions

- To unify task planning and motion planning layers, introduce notion of “domain semantics.”
- Domain semantics define how task operators map to motion planning problems.
- In TMKit, specified by Lisp or Python functions.
- TMKit also supports a meta-operator that “reparents” by modifying the underlying scene graph.



Contributions - IDTMP

- Iteratively Deepened Task and Motion Planning.
- “Secret sauce” of TMKit.
- Essentially:
 - Generate candidate task plans with set “step horizon”
 - Evaluate task plan feasibility using motion planner with set “sampling horizon”
 - If satisfying plan found, return
 - Else record found motion constraint or increment step horizon and sampling horizon, and loop.

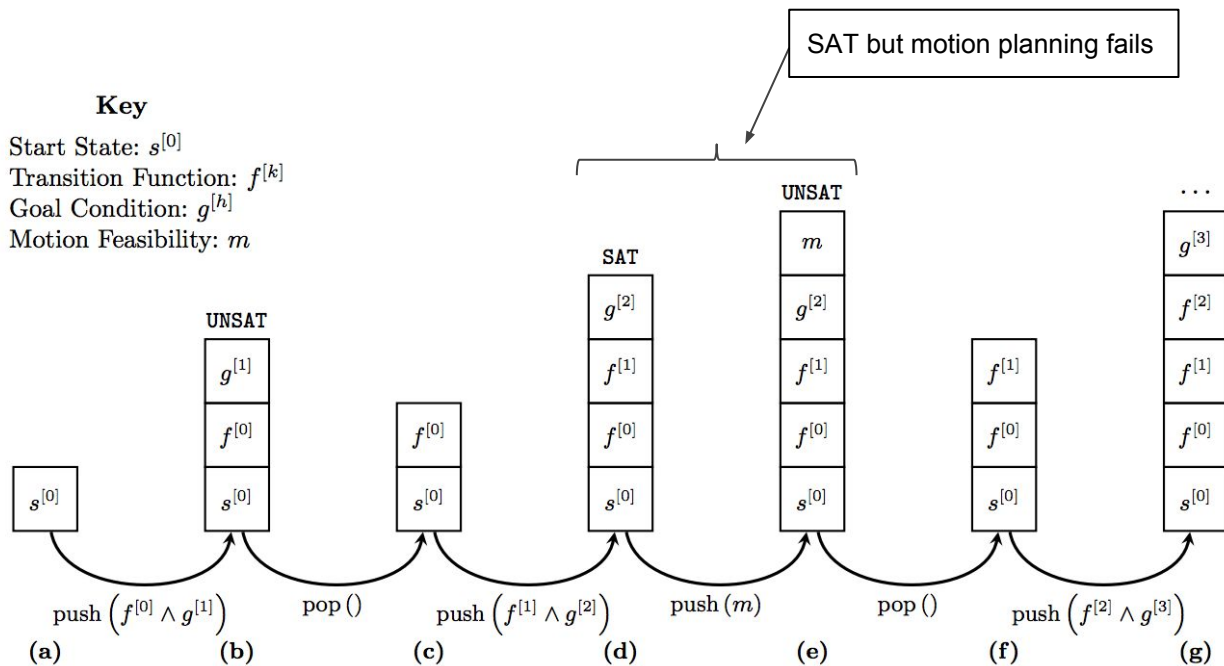


Contributions - IDTMP

- Incremental solving with SMT solver using a stack of “scopes” containing constraints.
- Constraints can be derived from geometric information returned by failed refinement of a task operator’s corresponding motion planning problem.
- The naive solution is to simply enumerate failed task plans as constraints.
- Better constraints developed in extensions.



Contributions - IDTMP



Contributions - Completeness

- An SMT solver will enumerate all candidate task plans at all step horizons, so if a feasible task plan exists in n steps, the solver is guaranteed to find it.
- RRT-Connect is probabilistically complete.
- Thus, since we retry failed motion plans by pushing and popping constraints, IDTMP is guaranteed to find a satisfying task and motion plan in the limit if one exists.



Contributions - Extensions to IDTMP


- By introducing post-plan and post-state connectivity assumptions, we can make use of more general constraints derived from failed motion plans.
- **Motion Plan Caching** - requires post-plan connectivity.
- **Failure Generalization Constraints** - requires post-state connectivity.
- **Collision Generalization Constraints** - requires post-state connectivity.
- All of these extensions preserve completeness since we pop the generated constraints eventually to deal with spurious motion planning failure.




Experiments - Scalability




(a)




(b)




(c)



(a)




(b)

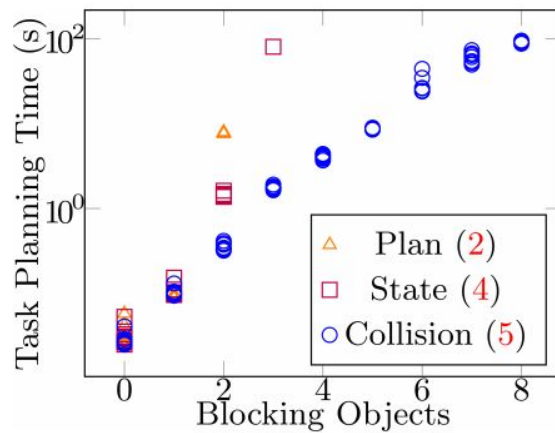


(c)


Experiments - Extensions



(a)



(b)



(c)

Dent

- Largely a technical accomplishment
 - Establishment of TMKit framework
 - Development of IDTMP algorithm
- Primary conceptual novelty in formalizations of task and motion planning concepts.
- Also exposes the inherent modularity of TMP in TMKit to facilitate future research.



Analysis - Literature Review

- Both papers make good use of interleaved references throughout.
- IJRR/RSS paper especially makes use of ~3 pages of references.
- Though task planning and combined task and motion planning are well covered by references, motion planning is less contextualized.
- Explanation of similar TMP approaches such as aSyMov and the Synergistic Framework is somewhat lacking in detail.



Analysis - Methodology

- One of the weaker points of the paper.
- Only three experiments performed, all in simulation, against only one alternate approach.
- Environments in which experiments are conducted pose no significant motion planning challenge.
- However, theoretical ideas are well justified, modularity seems reasonable for TMP, and IDTMP seems eminently extensible.
- Limitations to geometric motion planning are also acknowledged.



Analysis - Presentation

- Overall papers are well written, figures are informative.
- Occasional spelling errors and sentence fragments during related work discussion are distracting.
- Implementation-specific details and more general framework capabilities are not always clearly separated, especially in the RAM paper.



Questions?

